

# 301AA - Advanced Programming

Lecturer: **Andrea Corradini**

[andrea@di.unipi.it](mailto:andrea@di.unipi.it)

<http://pages.di.unipi.it/corradini/>

**AP-04:**      *Runtime Systems and intro to JVM*

# Overview

- Runtime Systems
- The Java Runtime Environment
- The JVM as an abstract machine
- JVM Data Types
- JVM Runtime Data Areas
- Multithreading
- Per-thread Data Areas
- Dynamic Linking
- JIT compilation
- Method Area

"HOW" THE PROGRAM IS EXECUTED

- TURING MACHINES
- LAMBDA CALCULUS
- ...
- != PARADIGM

# Runtime system

- Every programming language defines an **execution model**
- A **runtime system** implements (part of) such execution model, providing support during the execution of corresponding programs
- **Runtime support** is needed both by *interpreted* and by *compiled* programs, even if typically less by the latter

TO THE ABSTRACT  
MACHINE

# Runtime system (2)

- The runtime system can be made of
  - Code in the executing program generated by the compiler
  - Code running in other threads/processes during program execution IN JAVA FOR EXAMPLE THE GARBAGE COLLECTOR
  - Language libraries EXTERNAL FUNCTIONALITIES OF OUR NB. MACHINE
  - Operating systems functionalities
  - The interpreter / virtual machine itself

SYSTEM CALLS ARE PART OF THE RUNTIME SYSTEM

- I/O CALLS

PART OF THE RUNTIME SYS.  
SINCE THEY IMPLEMENT THE EXECUTION MODEL

# Runtime Support needed for...

- Memory management
  - Stack management: Push/pop of activation records
  - Heap management: allocation, garbage collection
- Input/Output
  - Interface to file system / network sockets / I/O devices
- Interaction with the **runtime environment**,
  - state values accessible during execution (eg. environment variables)
  - active entities like disk drivers and people via keyboards.

ALL THOSE DATA NEEDED BY THE PROGRAM IN EXECUTION  
BUT NOT PART OF THE PROGRAM ITSELF (OS BUFFERS, ...)

# Runtime Support needed for... (2)

- Parallel execution via threads/tasks/processes
- Dynamic type checking and dynamic binding
- Dynamic loading and linking of modules
- Debugging
- Code generation (for JIT compilation) and Optimization
- Verification and monitoring

COMPILED DURING EXECUTION

OF PROPERTIES THAT WE WANT TO HAVE  
DURING THE EXECUTION (EX. NO MEM. LEAK)  
REQUIRES RS AND MONITORING

TO DEBUG YOU NEED THE CURRENT VALUES OF VARS, BREAKPOINTS  
AND OTHER INFO AVAILABLE ONLY AT RUNTIME => NEED RS

# Java Runtime Environment - JRE

- Includes all what is needed to run compiled Java programs (COMPILED JAVA IS BYTESCODE)
  - JVM – Java Virtual Machine (THAT HAS AS MACHINE LANGUAGE THE BYTESCODE)
  - JCL – Java Class Library (Java API) (JAVA UTILS FOR EXAMPLE)
- We shall focus on the JVM as a real runtime system covering most of the functionalities just listed

- JVM UPDATES
- SIDE NOTE: JAVA VERSIONS CHANGE PRETTY OFTEN (WAY MORE THAN USED TO BE) BUT THEY MUST GUARANTEE RETROCOMPATIBILITY
    - EVERYTHING THAT WAS RUNNING ON THE OLD VERSION HAS TO RUN TO THE NEW VERSION
      - NOT EVERYONE DO THIS (EX: PYTHON 2 VS 3)

IMPOSE (ABSTRACT) PROPERTIES  
THAT HAVE TO BE RESPECTED,  
TRYING TO GIVE ZERO "HINTS"  
ABOUT THE IMPLEMENTATION

# What is the JVM?

ORACLE DONT IMPOSE AN  
IMPLEMENTATION OF THE JVM

- The **JVM** is an **abstract machine** in the true sense of the word.
- The **JVM specification** does *not* give implementation details like memory layout of run-time data area, garbage-collection algorithm, internal optimization (can be dependent on target OS/platform, performance requirements, etc.)
- The **JVM specification** defines a machine independent “**class file format**” that all **JVM implementations** must support
- The **JVM** imposes **strong syntactic and structural constraints** on the code in a class file. Any language with functionality that can be expressed in terms of a valid class file can be hosted by the **JVM**

ALLOWS STATIC  
ANALYSIS, SEE  
LATER ON

THE JVM IS NOT PORTABLE! BY NOT GIVING RULES OF IMPLEMENTATION BUT ONLY CONCEPTUAL PROPERTIES/CONSTRAINTS ANYONE CAN IMPLEMENT THE JVM TAILORED FOR HIS ARCHITECTURE

MORE THAN ONE FLOW OF EXECUTION  
GOING ON AT THE SAME TIME

- THREADS SHARE MEMORY
- GLOBAL, HEAP AND CODE

# Execution model

NOT RELATED TO  
THE CALL STACK

- JVM is a ***multi-threaded stack based machine***
- JVM instructions
  - implicitly take arguments from the top of the **operand stack** of the current frame
  - put their result on the top of the operand stack
- The operand stack is used to
  - pass arguments to methods
  - return a result from a method
  - store intermediate results while evaluating expressions
  - store local variables

WE WILL SEE LATER ON ↗

STACK BASED ARCHITECTURE: THE OP OF THE JVM OPERATES ONLY ON THE VALUES THAT ARE CONTAINED IN THE STACK. DATA HAVE TO BE BROUGHT IN THE STACK TO BE ABLE TO USE IT  
→ KINDA REGISTER MACHINE BUT WITH THE OPERAND STACK

- OPERATION READ/RP OPERANDS ON THE STACK AND PUSH ON TOP THE RESULT

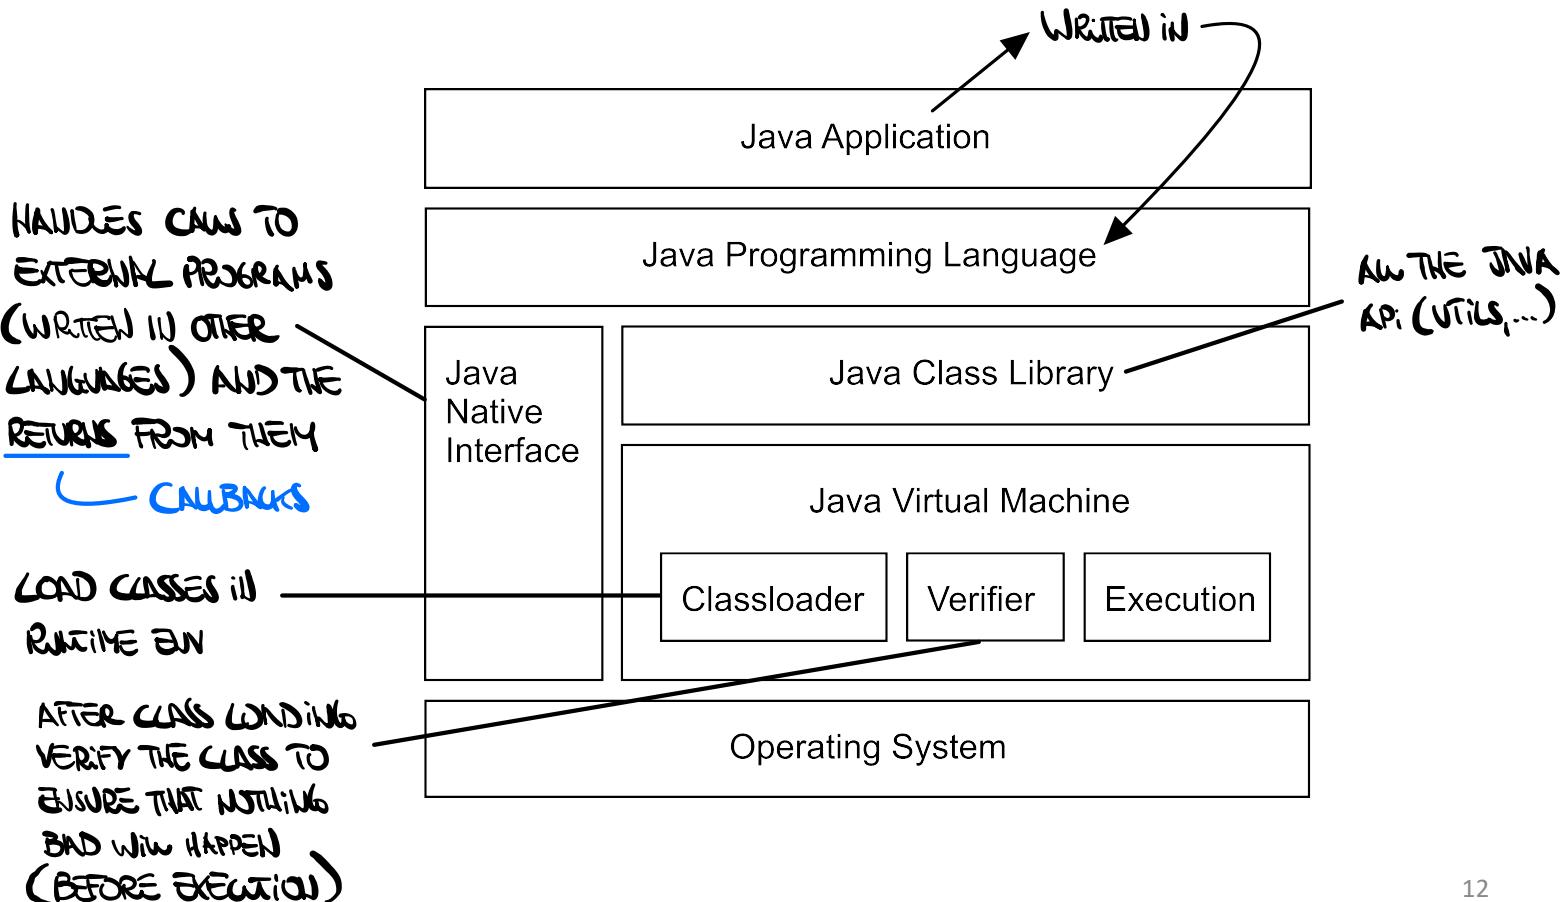
THE OPERAND STACK IS INSIDE THE ACTIVATION RECORD

CALL STACK

- EACH FUNCTION GENERATES A PUSH OF THE AR ON THE STACK AND INSIDE EACH AR THERE IS AN OPERAND STACK ON WHICH THE JVM DO THE STATE OF THE FUNCTION

⊕ HINTS DU LITTER

# Java Abstract Machine Hierarchy



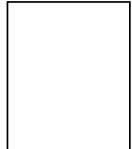
# Class Files and Class File Format

THE .CLASS FILES HAVE A PRECISE SYNTAX WHICH IS INDEPENDENT FROM THE ARCHITECTURE AND THE IMPLEMENTATION OF THE JVM

HOW STUFF IS IN MEMORY AND SO ON

External representation  
(platform independent)

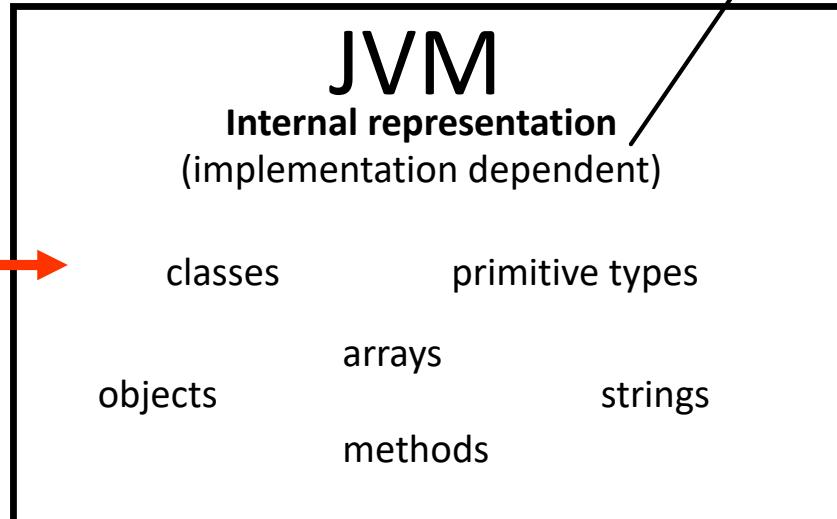
.class files



COMPILED JAVA CODE

JAVAC FILE.JAVA → FILE.CLASS

this is bytecode!



THE LOADING PROCESS TRANSFORMS CLASS FILES TO AN INTERNAL REPRESENTATION THAT IS SPECIFIC FOR THIS IMPLEMENTATION OF THE JVM

- THE LOADING PROCESS MUST BE IMPLEMENTED WITH THE IMPLEMENTATION OF THE JVM

Dimension defined by the  
JVM Spec, not implementation  
dependent!

# JVM Data Types

Unsigned short, 2 bytes  
as Unicode (not ASCII)

## Primitive types:

- numeric integral: byte, short, int, long, char \*
- numeric floating point: float, double, respectively 4 AND 8 BYTES
- boolean: boolean (support only for arrays) \*
- internal, for exception handling: returnAddress {CAN BE USED BY THE USER (DEV)}

## Reference types: (pointer (to heap) types: always 4 bytes)

- class types
- array types
- interface types

## Note:

- No type information on local variables at runtime \*
- Types of operands specified by **opcodes** (eg: iadd, fadd, ....)
- STRINGS ARE IMMUTABLE OBJECTS
- ARRAYS ARE OBJECTS BECAUSE THEY HAVE A TYPE-LIKE SYNTAX. SUPERCLASS OBJECT AND LOCATED ON THE HEAP

ADD ON INTEGERS

ADD ON FLOATS

- A LOT OF OPERATIONS ARE SUPPORTED ONLY ON INT, FLOAT, DOUBLE
  - TO SUM 2 BYTES I NEED TO CONVERT THEM TO INTEGERS  
THEN USE "iadd" AND CONVERT BACK THE RESULT
- BOOLEAN IS PRIMITIVE FOR THE LANGUAGE BUT NOT FOR THE JVM!  
OPERATIONS ARE NOT FULLY SUPPORTED

This is because the smallest indivisible word is of 1 byte hence we can't have operations on bool that target only 1 bit (can't read a bit, at least 1 byte)

  - BOOLEAN ARRAYS ( $> 8$  ELEMENTS) ARE SUPPORTED SINCE MULTIPLE CONCURRENTLY
  - TO SAVE A BIT (BOOLEAN) WE TAKE A FULL BYTE :C
- JVM OPERATIONS ARE TYPED AND THE RIGHT OPERATION IS CHOSEN USING THE TYPE ANNOTATIONS GIVEN BY THE USER (DEEWRITER)

# Object Representation IN THE MEMORY

- Left to the implementation (NOT DEFINED IN THE SPECIFICATION)
  - Including concrete value of null
- Extra level of indirection
  - need pointers to instance data and class data
  - make garbage collection easier
- Object representation must include
  - mutex lock
  - GC state (flags)

THE ACTUAL VALUE OF NULL  
IS NOT SPECIFIED, IT DEPENDS  
ON THE IMPLEMENTATION

- Most often is zero
- Not good choice, we'll see

IMPLEMENTATIONAL  
DEPENDENT

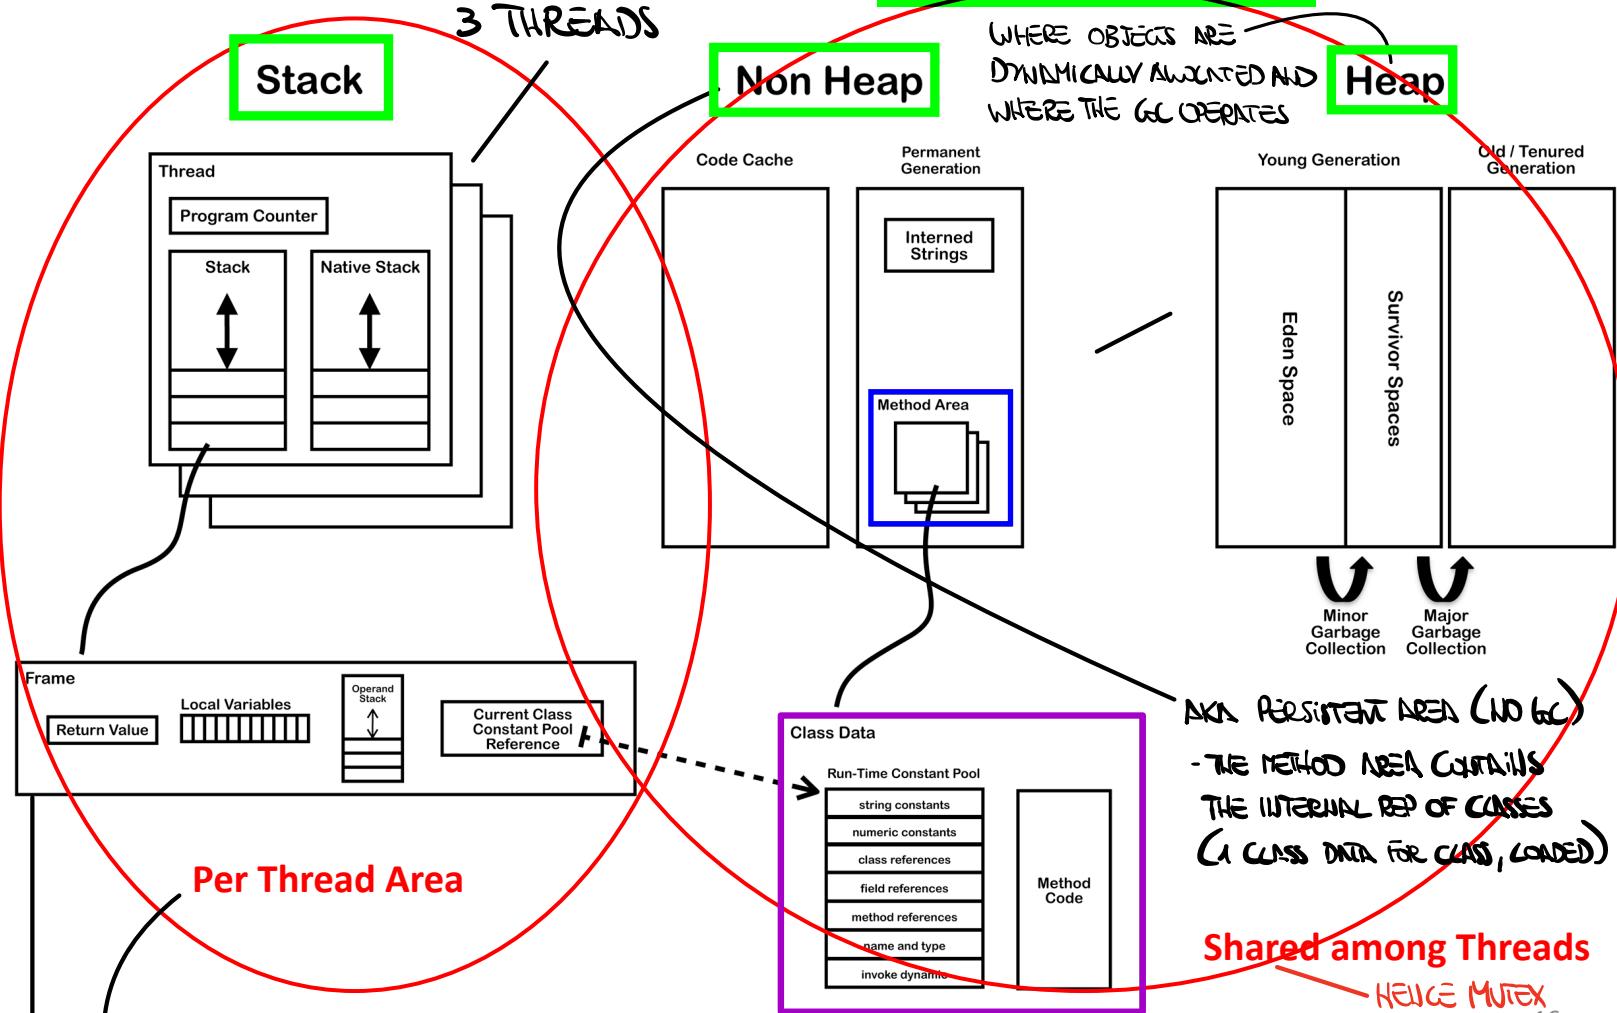
ALSO ALWAYS W TO DO  
REFERENCE COUNTING FOR  
GARBAGE COLLECTION

SINCE OBJECT STRUCTURE IS NOT DEFINED (IMPM. DEPEND.) WE CAN ALLOCATE OBJECTS  
ON THE STACK. ON THE STACK WE HAVE POINTERS (ALWAYS 4 BYTES) TO HEAP

- JVM IS CONCURRENT AND WITH GARBAGE COLLECTION THE JVM SPECIFIC IMPROVE THAT HAVE TO BE PRESENT MUTEX TO GUARANTEE MUTUAL EXCLUSION (IS THIS OBJECT LOCKED? WHICH THREAD HAS THE LOCK?)
-

# JVM Runtime Data Areas

THE FOLLOWING STUFF IS NEEDED IN THE JVM RDM SPECIFICATION



ONE THREAD AREA FOR EACH THREAD IN EXECUTION

- EACH THREAD HAS ITS OWN PROGRAM COUNTER THAT POINTS TO THE ADDRESS OF THE NEXT (BYTECODE) INSTRUCTION OF THE METHOD IN "EXECUTION" (WHICH IS FOUND BY THE CURRENT CLASS CONSTANT POOL REFERENCE), ITS OWN STACK OF AR TO HANDLE FUNCTIONS AND ITS NATIVE STACK, TO EXECUTE NATIVE CODE

IN EACH ACTIVATION RECORD WE HAVE

- OPERAND STACK: THE STACK USED BY THE STACK MACHINE (JVM) TO PERFORM THE BODY OF THE FUNCTION
- ARRAY OF LOCAL VARIABLES: THE LOCAL VARS USED BY THE FUNCTION RELATIVE TO THE AR
- RETURN VALUE: IT MAYBE COPED LATER ON BY THE JVM IF IT'S NEEDED TO ASSIGN THE VALUE TO SOMETHING ( $x = \text{OBJ.METHOD}()$ )
- CURRENT CLASS CONSTANT POOL REFERENCE: POINTS TO THE CLASS DATA THAT CONTAINS THE METHOD THAT WE ARE EXECUTING

- ◆ CLASS DATA CONTAINS THE RUNTIME CONSTANT POOL (NEED FOR THE DYNAMIC BINDING BETWEEN FUNCTION AND IMPLEMENTATION) AND THE NATIVE BYTECODE OF CLASS'S METHODS, TYPE OF PARAMS AND TYPE OF RETURNS

JVM IS MULTITHREADED AND HAS MULTIPLE  
THREADS IN EXECUTION FOR THE RUNTIME SUPPORT

- EVEN IF THE "USER" CODE IS SINGLE THREAD  
THERE ARE ALWAYS MORE THREADS RUNNING

# Threads

- JVM allows multiple threads per application, starting with `main` ————— 1ST THREAD STARTED
- Created as instances of `Thread` invoking `start()` (which invokes `run()`)
- Several background (daemon) system threads for
  - Garbage collection, finalization
  - Signal dispatching
  - Compilation, etc.

NOT PART OF THE APPLICATION LIFE CYCLE
- Threads can be supported by time-slicing and/or multiple processors (DEPENDS ON HW ARCHITECTURE AND UNDERLYING OS)

START DO THE INTERNAL ALLOCATION AND THEN INVOKE `run`, WHICH IS THE FUNCTION (PASSED AS PARAM)  
- THE THREAD WILL START A PARALLEL EXECUTION OF ITS CODE (`RUN METHOD CODE`)

# Threads (2)

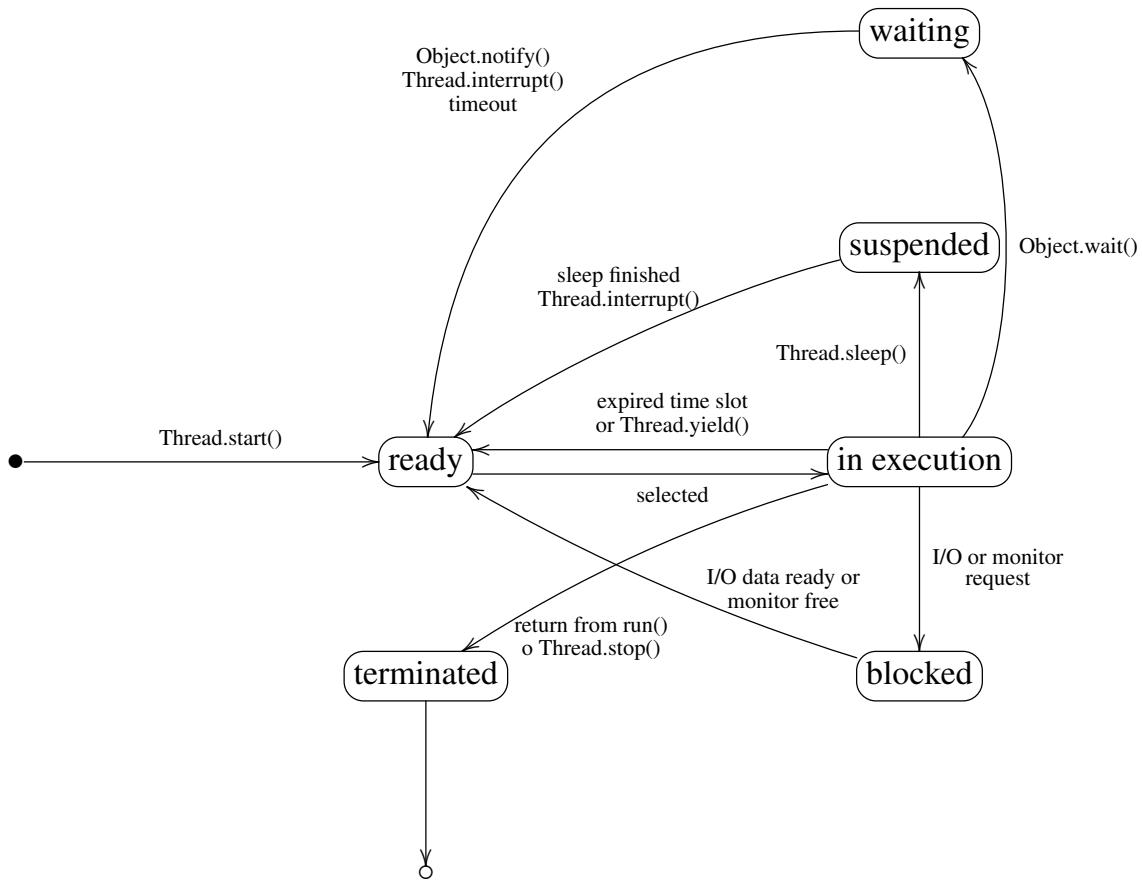
- Threads have shared access to heap and persistent memory
- Complex specification of consistency model
  - volatiles
  - working memory vs. general store
  - non-atomic longs and doubles

SPECIFY WHAT CAN HAPPEN  
WHEN THREAD CONCURRENTLY  
ACCESS A RESOURCE

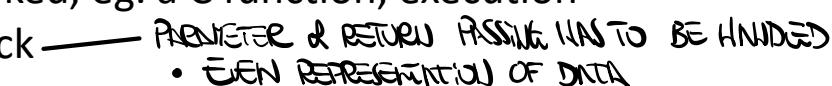
- THE SPECIFICATION OF THE CONSISTENCY MODEL HAS TO BE AS GENERAL AS POSSIBLE, NOT IMPOSING CONSTRAINTS OR IMPLEMENTATION
- EX: CANT IMPOSE 64 BIT ATOMIC READ AS MOST MACHINE WOULD SUPPORT IT

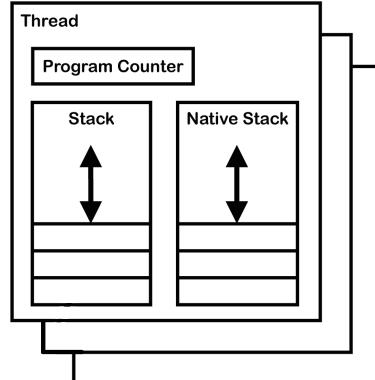
EXAMPLE: THE READ OF 32 BIT IS ATOMIC, THE READ OF 64 BIT IS DONE WITH 2 ATOMIC READS: WHAT HAPPENS IF SOMEONE DOWNSIZES BETWEEN READS?

# Java Thread Life Cycle



# Per Thread Data Areas

- pc: pointer to next instruction in *method area* 
  - undefined if current method is *native* 
- The **java stack**: a stack of *frames* (or *activation records*).
  - A new frame is created each time a method is invoked and it is destroyed when the method completes.
  - The JVMS does not require that frames are allocated contiguously
- The **native stack**: is used for invocation of native functions, through the JNI (Java Native Interface)
  - When a native function is invoked, eg. a C function, execution continues using the native stack 
    - EVEN REPRESENTATION OF DATA
  - Native functions can call back Java methods, which use the Java stack



WE REPRESENT IT CONTINUOUSLY BUT THAT IS NOT IMPOSED: COULD BE IMPLEMENTED AS A LINKED LIST (THE JVM IS AN ABSTRACT MACHINE)

④ PC POINTS TO THE NEXT LINE OF CODE TO EXECUTE, HOW WE FIND IT?

- IN THE ("CURRENT") AR OR IN THE STACK WE HAVE THE POINTER TO THE CLASS DATA  
(THE POINTER IS ON THE METHOD AREA) 

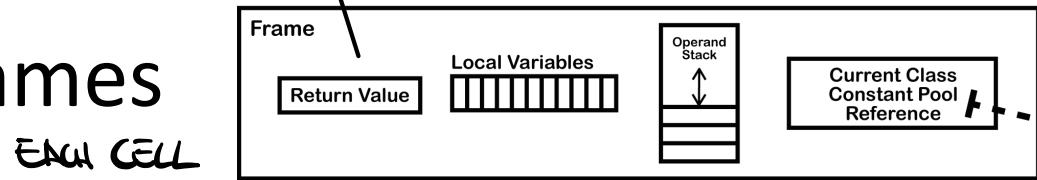
- THE CLASS DATA CONTAINS THE METHOD CODE (THE ACTUAL BYTECODE)

⑤ BY SPECIFICATION IF THE CODE IN EXECUTION IS NOT BYTECODE (EX: C THRU JNI) THE JVM DO  
NOT EVEN KNOW THE EXECUTION MODEL OF THAT (NATIVE) CODE => PC UNDEFINED

⑥ (SEE BELOW) THE SIZE OF BOTH IS NOT FIXED BUT ESTIMATED AT COMPILE TIME  
(WHEN THE BYTECODE IS GENERATED): VERY EFFICIENT BECAUSE FOR SMALL METHODS  
WE DO NOT WASTE MEMORY

VALUE OF RETURN OF THE FUNCTION (IF ANY). AT THE  
DESTRUCTION OF THE AR (WHEN THE FUNCTION RETURNS)  
THE VALUE IS COPIED ON THE LOCAL VARIABLES OF THE  
CALLER (IF THE RETURN VALUE IS ASSIGNED TO SOMETHING)

# Structure of frames



- **Local Variable Array (32 bits)** containing
  - Reference to this (if instance method) *in Pos 0 of THE ARRAY*
  - Method parameters
  - Local variables
- **Operand Stack** to support evaluation of expressions and evaluation of the method
  - Most JVM bytecodes manipulate the stack
- Reference to **Constant Pool** of current class
  - EACH METHOD IS DEFINED IN SOME CLASS (STATIC OR INSTANCE OF..);  
- THIS POINTS TO THE (INTERNAL) REP OF THE CURRENT CLASS OR CONST. POOL THAT CONTAINS THE CODE (BYTECODE)

SUBSTITUTION SYMBOLIC ADDRESSES OF ALREADY COMPILED CODE WITH  
ACTUAL PHYSICAL ADDRESSES THAT CAN BE USED IN THE EXECUTION

# Dynamic Linking (1)

MIGHT BE API, LIBRARIES...

- The reference to the constant pool for the current class helps to support **dynamic linking**.
- In C/C++ typically multiple object files are linked together to produce an executable or dll.
  - During the linking phase symbolic references are replaced with an actual memory address relative to the final executable.
- In Java this linking phase is done **dynamically** at runtime.
- When a Java class is compiled, all references to variables and methods **are stored in the class's constant pool as symbolic references**.

IN THE BYTECODE THERE ARE LOGIC ADDRESSES!  
THEY ARE RESOLVED ONLY AT RUNTIME DURING THE EXECUTION

# Dynamic Linking (2)

- The JVM implementation can choose when to resolve symbolic references.
  - **Eager or static resolution:** when the class file is verified after being loaded (✿)
  - **Lazy or late resolution:** when the symbolic reference is used for the first time (✿)
- The JVM **has to behave** as if the resolution occurred when each reference is first used and throw any resolution errors at this point. { (✿)
- **Binding** is the process of the entity (field, method or class) identified by the symbolic reference being replaced by a direct reference
- This only happens **once** because the symbolic reference **is completely replaced** in the constant pool
- If the symbolic reference refers to a class that has not yet been resolved then this class will be loaded. IN THE METHOD AREA AND THEN RESOLVED AFTER THE 1ST TRANSACTION ; REPLACE THE FREE OCCURENCE OF THE LOGIC ADDRESS

## • LOADING PHASE THEN VERIFICATION PHASE

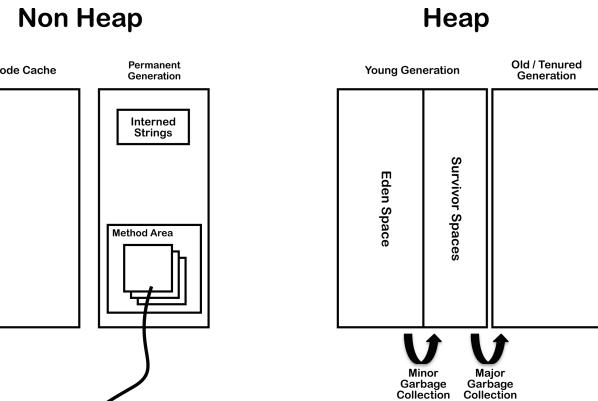
- IN EAGER RESOLUTION IF IN THE VERIFICATION PHASE IS FOUND AN UNRESOLVED REFERENCE WE RESOLVE IT
  - THE CLASS THAT CONTAINS THE DATA (EG. METHOD ATTRIBUTE) AND THE LOGIC NODNESS GET TRANSLATED
    - EAGER: WE RESOLVE ADDRESSES THAT MAY BE NEVER USED IN THE EXECUTION (AN IF BRANCH...)
- THE EAGER RESOLUTION MAKES THE CLASS CONTAIN A STATIC ENV (DOE STATICALLY, NOT AT RUNTIME)
  - THE EXECUTION WILL BE FASTER BUT ERRORS MIGHT BE SIGNIFIED EVEN IF THEY WOULD NOT OCCUR IN THE EXECUTION

• GET'S RESOLVED THE FIRST TIME USED AND STORED (NEVER TRANSLATE ADDRESSES TWICE)

• THE SPECIFICATION STATES THAT THE JVM HAS TO BEHAVE AS THE RESOLUTION IS DONE WHEN THE REFERENCE IS FIRST USED (AND THROWN ERRORS IF NOT)

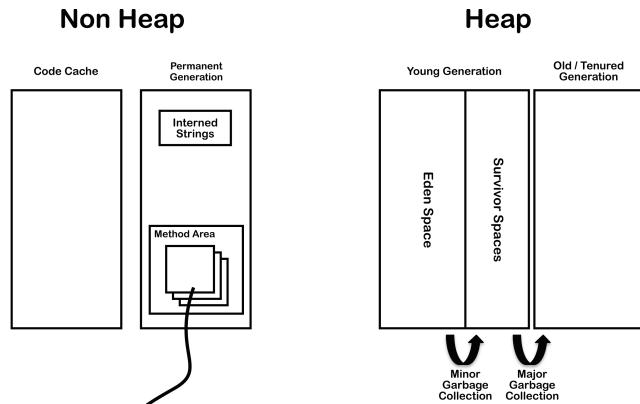
→ IF YOU USE AN EAGER APPROACH YOU MAY ALREADY KNOW FROM THE LOADING PHASE THAT THERE IS A PROBLEM, BUT YOU CAN'T REPORT IT UNTIL THE REFERENCE IS USED

# Data Areas Shared by Threads: Heap



- Memory for objects and arrays; unlike C/C++ they are never allocated to stack
- Explicit deallocation not supported. Only by garbage collection.
- Since Oracle JDK 11: **Z Garbage Collector**

# Data Areas Shared by Threads: **Non-Heap**



- Memory for objects which are never deallocated, needed for the JVM execution
  - Method area — HERE THE INTERNAL REP OF CLASSES
  - Interned strings — EXPLICIT STRINGS BETWEEN " ". STILL OBJECTS AND THEY ARE STORED HERE SINCE THEY ARE NOT ALLOCATED MULTIPLE TIMES
    - STRING IS FINAL AND STRINGS ARE IMMUTABLE
  - Code cache for JIT

GENERALLY IN THE EXECUTION OF A PROGRAM THERE IS A PART OF THE MEMORY THAT IS STATIC (LIKE THE CODE) THAT IS ALLOCATED WHEN THE PROGRAM STARTS AND DEALLOCATED WHEN THE PROGRAM ENDS. THEN THERE'S THE DYNAMIC PART OF THE MEMORY LIKE THE HEAP THAT GROWS AND SHRINKS DURING THE EXECUTION

- THIS **NON-HEAP** IS SEMI STATIC: NOT FIXED COMPLETELY AT THE START (METHOD AREA FOR EXAMPLE MAY GROW) BUT IS ALMOST NEVER REFERENCED BY THE GC

# JIT (Just In Time) compilation

ORACLE'S IMPLEMENTATION OF THE JVM

- The **Hotspot JVM** (and other JVMs) profiles the code during interpretation, looking for “hot” areas of byte code that are executed regularly
- These parts are compiled to native code.
- Such code is then stored in the **code cache** in non-heap memory.
- EFFICIENCY BOOST IN EXECUTION TIME

# Method area

The memory where class files are loaded. For each class:

- Classloader Reference
- From the class file:
  - Run Time Constant Pool
  - Field data (w Scoping, static,..)
  - Method data
  - Method code (Bytecode)

WE WILL SEE IT LATER ON

REMEMBER: .CLASS FILES ARE IMPLM.  
INDEPENDENT! DURING THE LOADING PHASE  
THERE IS A PASSAGE FROM MACHINE INDP.  
TO INTERNAL REPRESENTATION

IN THE .CLASS FILE THERE IS THE STATIC  
CONSTANT POOL, THAT IN THE LOADING IS  
TRANSFORMED "1-TO-1" IN THE RT CONST P

**Note:** Method area is shared among thread. Access to it has to be  
**thread safe.**

Changes of method area when:

- A new class is loaded
- A symbolic link is resolved by dynamic linking

THE RESOLUTION IS DONE BY  
REPLACING THE LOGICAL ADDRESSES  
WITH THE PHYSICAL ADDRESSES IN  
THE RUNTIME CONSTANT POOL  
- HENCE THIS PROCESS AFFECTS  
THE METHOD AREA

FOR EACH METHOD:

# ARGUMENTS, ARG TYPES, RETURN TYPE

UH: UNSIGNED IN BYTES

# Class file structure (BY JVM SPECIFICATION)

ClassFile {

u4 magic;	0xCAFEBABE
u2 minor_version;	Java Language Version
u2 major_version;	
u2 constant_pool_count; — #ENTRIES IN STATIC CONST POOL	Constant Pool
cp_info contant_pool[constant_pool_count-1]; ACTUAL CONST POOL	
u2 access_flags; (PRIVATE, PUBLIC, PROTECTED,...)	access modifiers and other info
u2 this_class;	References to Class and Superclass
u2 super_class;	
u2 interfaces_count; #INTERFACES OF THIS CLASS	References to Direct Interfaces
u2 interfaces[interfaces_count];	
u2 fields_count; #FIELDS	Static and Instance Variables
field_info fields[fields_count];	ACTUAL FIELDS
u2 methods_count; #METHODS	Methods
method_info methods[methods_count];	METHODS CODES (BYTECODE)
u2 attributes_count;	
attribute_info attributes[attributes_count];	Other Info on the Class

{METHODS ACCESS (PRIVATE, PUBLIC,...),  
SIGNATURES, #NRS & TYPES, RETURN TYPE}

DATA CARE

# Field data in the Method Area

Per field:

- Name
- Type
- Modifiers
- Attributes

*FIELD\_info IN THE CLASS FILE*

USED FOR THE RESOLUTION OF SYMBOLIC ADDRESSES  
SINCE THE TRANSLATION IS DYNAMIC

- STATIC RESOLUTION LANGUAGES CAN FORGET  
NAMES AND RESEND ONLY WITH PHYSICAL  
ADDRESSES

PUBLIC, PRIVATE, ..

/  
FOR ANNOTATIONS

# FieldType descriptors

INSIDE CLASSFILE

<i>FieldType</i> term	Type	Interpretation
B	byte	signed byte
C	char	Unicode character code point in the Basic Multilingual Plane, encoded with UTF-16
D	double	double-precision floating-point value
F	float	single-precision floating-point value
I	int	integer
J	long	long integer
L <i>ClassName</i> ;	reference	an instance of class <i>ClassName</i>
S	short	signed short
Z	boolean	true or false
[	reference	one array dimension

THEN: [I MATRIX, D CUBE, ...<sup>30</sup>

# Method data

Per method:

- Name
- Return Type
- Parameter Types (in order)
- Modifiers
- Attributes
- Method code...

A *method descriptor* contains

- a sequence of zero or more parameter descriptors in brackets
- a return descriptor or V for void descriptor

Example: The descriptor of

Object m(int i, double d, Thread t) {...}

is:

(IDLjava/lang/Thread;)Ljava/lang/Object;

# Method code

ALL THIS INFO ARE NEEDED  
FOR METHOD EXECUTION!

Per method:

- Bytecodes
  - Operand stack size }  
• Local variable size }  
• Local variable table }  
• Exception table ↗  
• LineNumberTable – which line of source code corresponds to  
which byte code instruction (for debugger)
- COMPILED ESTIMATION OF SIZE OF THE  
OPERAND STACK AND LOCAL VARIABLES ARRAY  
LATER ON

Per exception handler (one for each try/catch/finally clause)

- Start point
- End point
- PC offset for handler code →  $PC\_TRY + PC\_OFFSET = PC\_CATCH$
- Constant pool index for exception class being caught

## \* THE EXCEPTION TABLE OF A METHOD SAYS

- PROGRAM COUNTER OF THE START AND THE END OF THE BLOCK "TRY" IN THE BYTCODE
  - REFERENCE TO THE EXCEPTION CLASSES THAT THE CATCH BLOCK(S) CAPTURES
  - PC OFFSET IN THE BYTCODE TO JUMP TO THE HANDLER
- 
- THE EXCEPTION HANDLING PROCEDURE IS AS FOLLOWS
    - IF AN EXCEPTION IS GENERATED BY AN INSTRUCTION (MAYBE A METHOD INVOCATION OR WHWR) FIND THE CORRESPONDING TRY INTERVAL, FIND THE CORRESPONDING EXCEPTION TYPE AND THEN JUMP TO THE CATCH BYTCODE USING THE PC OFFSET SAW EARLIER

# Disassembling Java files: javac, javap, java

`SimpleClass.java`

```
package org.jvminternals;  
public class SimpleClass {  
    public void sayHello() {  
        System.out.println("Hello");  
    }  
}
```



Compiler  
`javac SimpleClass.java`

BINARY  
SimpleClass.class

BINARY  
INTERPRETER

JVM  
`java SimpleClass`

TO GET A READABLE VERSION OF BINARY  
- PROCESS CALLED DISASSEMBLING

Disassembler  
`javap -c -v SimpleClass.class`

# SimpleClass.class: constructor and method

THERE SHOULD BE THE CONSTANT POOL: Skipped next slide

Local variable 0 = "this"

```
{  
public org.jvminternals.SimpleClass();  
descriptor: ()V  
flags: ACC_PUBLIC  
Code:  
stack=1, locals=1, args_size=1  
0: aload_0  
1: invokespecial #1 // Method java/lang/Object."<init>": ()V  
4: return  
LineNumberTable:  
line 2: 0
```

```
package org.jvminternals;  
public class SimpleClass {  
    public void sayHello() {  
        System.out.println("Hello");  
    } }
```

```
public void sayHello();  
descriptor: ()V  
flags: ACC_PUBLIC  
Code:  
stack=2, locals=1, args_size=1  
0: getstatic      #2      // Field java/lang/System.out:Ljava/io/PrintStream;  
3: ldc            #3      // String Hello  
5: invokevirtual #4      // Method  
java/io/PrintStream.println:(Ljava/lang/String;)V  
8: return  
LineNumberTable:  
line 4: 0  
line 5: 8  
}  
SourceFile: "SimpleClass.java"
```

Index into constant pool

Method descriptors

THE DISASSEMBLER RIS IN  
CONSTANTS FOR READABILITY

String literal

Field descriptor

## ⌚ INFO ABOUT CLASS DESCRIPTOR

- METHOD DESCRIPTOR: ()V : NO PARAMS (), RETURNS VOID V
- FLAGS: ACC\_PUBLIC, THE METHOD IS PUBLIC
- CODE
  - STACK: ESTIMATE SIZE OF THE OPERAND STACK (WHERE BYTECODE IS EXECUTED BY THE JVM)
  - LVAL: LOCAL VARIABLE ARRAY SIZE GUESS (1, BECAUSE  $\$ = \text{THIS}$ )
  - ARGS\_SIZE: # ARGUMENTS PASSED TO THE METHOD (1, JUST THIS)
- UNDER THE INFO THERE IS THE ACTUAL BYTECODE
  - ALLOC\_0: PUSH ON THE STACK THE ELEMENT IN POS 0 OF THE LOCAL VARIABLE ARRAY
    - THE ANALOGOUS OPERATION TO SAVE FROM STACK TO LOCAL VAR ARRAY IS STORE
  - INVOKESTATIC #1 : INVOKING INSTANCE METHOD ; SPECIAL HANDLING FOR SUPERCLASS, PRIVATE, ... . #1 IS A SYMBOLIC REFERENCE TO THE CONSTANT POOL, WHEN RESOLVED IT WILL BE THE ADDRESS OF THE CONSTRUCTOR OF THE (SUPER)CLASS OBJECT
- LINENUMBERTABLE: DON'T CARE

## ⌚ SAME REASONING



IN COMPUTERS THE SYMBOL TABLE  
MATCH TO EACH VARIABLE NAME  
SOME INFO (MEM. LOCATION, TYPE, ...)

# The constant pool

THESE ARE THE TYPES THAT WE  
CAN FIND IN THE CONSTANT POOL  
OF THE .CLASS

- Similar to *symbol table*, but with more info
- Contains constants and symbolic references used for dynamic binding, suitably tagged
  - numeric literals (Integer, Float, Long, Double)
  - string literals (Utf8)
  - class references (Class)
  - field references (Fieldref)
  - method references (Mehodref, InterfaceMethodref, MethodHandle)
  - signatures (NameAndType)
- Operands in bytecodes often are indexes in the constant pool

OPJ IN THE STACK USES LOCAL #INDEXES THAT, FOLLOWING A PATH  
ARE RESOLVED INTO PHYSICAL ADDRESS (SEE NEXT SLIDE)

STATIC CONSTANT POOL RELATED BY THE COMPILED OF SimpleClass.java

# SimpleClass.class: the Constant pool

Compiled from "SimpleClass.java"  
public class SimpleClass  
minor version: 0  
major version: 52  
flags: ACC\_PUBLIC, ACC\_SUPER

Constant pool:

#1 = Methodref	#6..#14
#2 = Fieldref	#15..#16
#3 = String	#17
#4 = Methodref	#18..#19
java/io/PrintStream.println:(Ljava/lang/String;)V	
#5 = Class	#20
#6 = Class	// SimpleClass
#7 = Utf8	#21
#8 = Utf8	<init>
#9 = Utf8	()V
#10 = Utf8	Code
#11 = Utf8	LineNumberTable
#12 = Utf8	sayHello
#13 = Utf8	SourceFile
#14 = NameAndType	SimpleClass.java
#15 = Class	#7:#8
#16 = NameAndType	// <init>: ()V
#17 = Utf8	#22
#18 = Class	// java/lang/System
#19 = NameAndType	#23:#24
#20 = Utf8	Hello
#21 = Utf8	#25
#22 = Utf8	// out:Ljava/io/PrintStream;
#23 = Utf8	// java/io/PrintStream
#24 = Utf8	#26:#27
#25 = Utf8	// println:(Ljava/lang/String;)V
#26 = Utf8	SimpleClass
#27 = Utf8	java/lang/Object
	java/lang/System
	out
	Ljava/io/PrintStream;
	java/io/PrintStream
	println
	(Ljava/lang/String;)V

```
public class SimpleClass {  
    public void sayHello() {  
        System.out.println("Hello");  
    } }
```

SOURCE CODE

// java/lang/Object.<init>: ()V  
// java/lang/System.out:Ljava/io/PrintStream;  
// Hello  
//  
<init>  
()V  
Code

ALL THESE ADDRESSES ARE DEFINED BY  
THE COMPILER SO THAT CAN BE RESOLVED

STRUCT TO PRINT DIRECTLY IN THE CLASS  
STRUCTURE

BYTECODE

```
public void sayHello();  
descriptor: ()V  
Code:  
stack=2, locals=1, args_size=1  
0: getstatic #2  
3: ldc #3  
5: invokevirtual #4  
8: return
```

• HERE METHODS INFO (Saw Before)

AT THIS POINT WE RESOLVE THE ADDRESS OF THE CLASS SYSTEM AND WE LOAD IN THE METHOD AREA THE CORRESPONDING INTERNAL REPRESENTATION (OF SYSTEM)

④ GETSTATIC TO PUSH IN THE STACK THE STATIC FIELD (`o1`) OF SYSTEM (AND ON (`o1` WE WILL INVOKE PRINTLN))

- IDC PUSH ON THE OPERAND STACK A REFERENCE CONTAINED IN THE CONSTANT POOL. IN THIS CASE THE REFERENCE POINTS TO THE STRING "Hello" (FOLLOW THE PATH: #3 → #17 → "Hello")

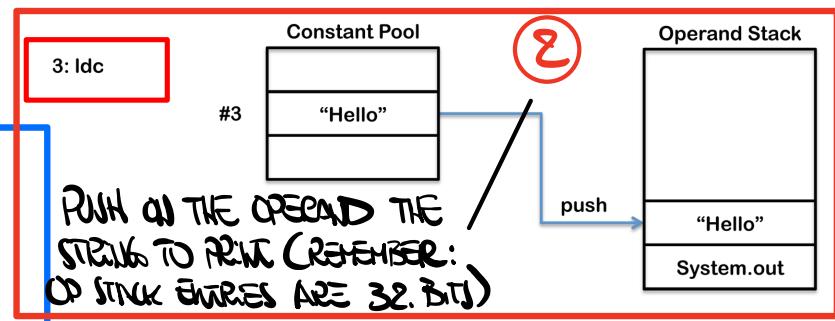
⑤ (SEE BELOW)

TO THE INVOKEVIRTUAL IT IS PASSED THE DESCRIPTOR OF THE FUNCTION PRINTLN, THE OPERAND OF INVOKEVIRTUAL IS TO CALL PRINTLN ON THE Z (FIRST) VALUE ON THE STACK (WE PUSHED THIS = `o1` AND THE PARAMETER "Hello")

- THE EXECUTION PROCEEDS BY CREATING A NEW ACTIVATION RECORD IN THE THREAD AREA OF THE CURRENT THREAD, WHERE IN ITS OPERAND STACK WE HAVE `o1` AND "Hello". WHEN THE METHOD PRINTLN ENDS (IT'S DE DESTROYED) WE WILL FIND THE RETURN VALUE ON THE STACK (IN DR) OF THE OWNER (SAY `Hello`) (PRINTLN IS VOID SO IN THIS CASE NOTHING)

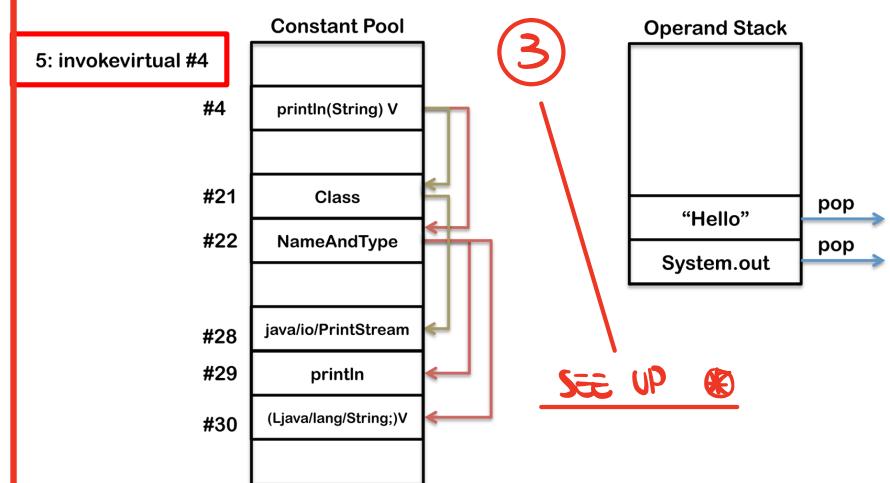
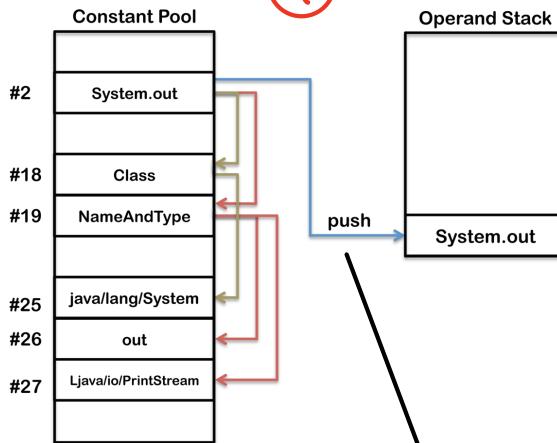
# BYTESIDE EXECUTION SIMULATOR

```
public void sayHello();
descriptor: ()V
Code:
stack=2, locals=1, args_size=1
0: getstatic      #2
3: ldc           #3
5: invokevirtual #4
8: return
```



sayHello()

0: getstatic



GETSTATIC PUSH ON THE OPERAND STACK  
THE REFERENCE OF OUT (STATIC FIELD OF SYSTEM)

# Loading, Linking, and Initializing

TRANSFORM .CLASS STATIC INFO (BINARY REP) IN INTERNAL REP INTO

- **Loading:** finding the binary representation of a class or interface type with a given name and creating a class or interface from it —  
THE CLASS MIGHT ARRIVE FROM WEB
- **Linking:** taking a class or interface and combining it into the run-time state of the Java Virtual Machine so that it can be executed
- **Initialization:** executing the class or interface initialization method <clinit>

| THERE COULD BE STATIC BLOCKS TO BE EXECUTED AT LOADING TIME  
| (MAYBE TO INITIALIZE FIELDS OR OTHER). THESE BLOCKS ARE CALLED CLINIT

CLASSIC PROCEEDING OF THE PROGRAM STARTS

## JVM Startup

MIND THAT THE POWER OF CLASS LOADING DEPENDS  
ON THE SPECIFIC IMPLEMENTATION OF THE JVM

- The JVM starts up by loading an initial class using the **bootstrap classloader**
  - The class is linked and initialized
  - `public static void main(String[])` is invoked.
  - This will trigger loading, linking and initialization of additional classes and interfaces...
- } THE BOOTSTRAP CLASSLOADER LOADS THE CLASS IN WHICH THE MAIN IS CONTAINED, IT GETS LINKED AND THEN ——————

① PARSE THE BINARY DATA FROM  
.CLASS FILE ACCORDING TO SPECIFICATION,  
THEN TRADE IT IN THE IMPLEMENTATION  
DEPENDENT INTERNAL REP

② CREATED AN INSTANCE OF  
JAVA. LANG. CLASS

# Loading

- Class or Interface C creation is triggered
  - by other class or interface referencing C
  - by certain methods (eg. reflection)
- Array classes are generated by the JVM
- Check whether already loaded
- If not, invoke the appropriate loader.loadClass
- Each class is tagged with the *initiating loader*

ARRAY CLASSES NOT  
LOADED BUT GENERATED  
ON USAGE BY THE JVM  
- GENERATE AND LOAD  
AT USAGE FOR EACH  
ARRAY  
- NEW CLASS IN  
CHANGE TYPE OR SIZE

THE REFERENCE IS IN THE CNTS. POOL

ALSO LOADS THE  
CLASS WHERE THERE'S  
THE MAIN METHOD

# Class Loader Hierarchy

- **Bootstrap Classloader** loads basic Java APIs, including for example `rt.jar`. It may skip much of the validation that gets done for normal classes.  
*BASIC JAVA APIs, ALREADY VALIDATED*
- **Extension Classloader** loads classes from standard Java extension APIs such as security extension functions.
- **System Classloader** is the default application classloader, which loads application classes from the classpath
- **User Defined Classloaders** can be used to load application classes:
  - for runtime reloading of classes — *A CLASS MAY CHANGE DYNAMICALLY AT RUNTIME*
  - for loading from different sources, eg. from network, from an encrypted file, or also generated on the fly
  - for supporting separation between different groups of loaded classes as required by web servers
- Class loader hooks: `findClass` (builds a byte array), `defineClass` (turns an array of bytes into a class object), `resolveClass` (links a class)

SEPARATE THE CLASS REP OF THE SAME CLASS FOR DIFFERENT USERS ACCESSING  
THE WEB SERVER (THEY USE THE SAME JVM, TO EXPENSIVE OTHERWISE)

FIND CLASS: READS THE BINARY OF THE CLASS (FROM THE FILE OR ONLINE)  
DEFINE CLASS: INTERPRETER THE BINARY TRANSFORMING IT IN A CLASS (INTO INTERNAL REP)  
RESOLVE CLASS: RESOLVE DYNAMIC REFERENCES

# Runtime Constant Pool

THE RUNTIME CONST POOL IS THE INTERNAL REP OF THE STATIC CONST POOL

- The constant\_pool table in the .class file is used to construct the *run-time constant pool* upon class or interface creation.
- All references in the run-time constant pool are initially symbolic. (THEY ARE RESOLVED UPON USE)
- Symbolic references are derived from the .class file in the expected way
- Class names are those returned by **Class.getName()** AS PER THE OBJ OF WHICH WE WANT THE CLASS
- Field and method references are made of name, descriptor and class name

A LINKER IS A PROGRAM THAT TAKES INDIVIDUAL COMPILED FILES AND COMBINES THEM INTO A SINGLE EXECUTABLE PROGRAM

# Linking

OPTIONAL: THE RESOLUTION CAN BE EAGER OR LAZY  
- IF DONE IN LINKING PHASE IS EAGER

- Link = verification, preparation, resolution
- **Verification:** see below
- **Preparation:** allocation of storage (method tables)
- **Resolution (optional):** resolve symbol references by loading referred classes/interfaces
  - Otherwise postponed till first use by an instruction

ONCE THE VERIFICATION IS DONE THE JVM ALLOCATES THE MEMORY FOR THE CLASS VARIABLES AND INITIALIZED THEM TO THE DEFAULT VALUE ACCORDING TO THE TYPE OF VARIABLE.

- THE ACTUAL INITIALIZATION (WITH USER DEFINED VALUES) DO NOT OCCUR UNTIL INITIALIZATION PHASE

# Verification

ALSO CANT KNOW IF  
ALREADY VERIFIED

- When?
  - Mainly during the load and link process
- Why?
  - No guarantee that the class file was generated by a Java compiler
  - Enhance runtime performance
- Examples
  - There are no operand stack overflows or underflows.
  - All local variable uses and stores are valid.
  - The arguments to all the JVM instructions are of valid types.
- Relevant part of the JVM specification: described in ~170 pages of the JVMS (total: ~600 pages)

AT RUNTIME

VERIFICATION HAS SAFE POWER : IF PASSED NO ERRORS, IF NOT PASSED  
THERE IS A CHANCE THAT STILL NOT ERROR

# Verification Process

ONLY A SYNTACTIC  
STEP: CHECK THAT THE  
.CLASS IS OK BY THE  
CLNS FILE SPECIFICATION

- Pass 1 – when the class file is loaded
    - The file is properly formatted, and all its data is recognized by the JVM
  - Pass 2 – when the class file is linked
    - All checks that do not involve instructions
      - final classes are not subclassed, final methods are not overridden.
      - Every class (except Object) has a superclass.
      - All field references and method references in the constant pool have valid names, valid classes, and a valid type descriptor.
- OTHER CHECKS  
- NO OVERFLOW OR UNDERFLOW  
- VARIABLES ARE INITIALIZED  
- TYPECHECKING ON METHODS BASED ON SIGNATURES
- SEMANTIC VERIFICATION STEP

# Verification Process – cont.

- Pass 3 – still during linking
  - **Data-flow analysis** on each method.
  - Ensure that at any given point in the program, no matter what code path is taken to reach that point:
    - The operand stack is always the same size and contains the same types of objects.
    - No local variable is accessed unless it is known to contain a value of an appropriate type.
    - Methods are invoked with the appropriate arguments.
    - Fields are assigned only using values of appropriate types.
    - All opcodes have appropriate type arguments on the operand stack and in the local variables
    - A method must not throw more exceptions than it admits
    - A method must end with a return value or throw instruction
    - Method must not use one half of a two word value

# Verification Process – cont.

- Pass 4 - the first time a method is actually invoked
  - a virtual pass whose checking is done by JVM instructions
  -  • The referenced method or field exists in the given class.
  - The currently executing method has access to the referenced method or field.

RUNTIME CHECKS

- BOTH CHECKS REQUIRE THE RESOLUTION OF DYNAMIC REFERENCE
  - DONE ON THE 1<sup>ST</sup> USAGE OF THE METHOD

(SYNTAX SPECIFIED BY THE JVM): STATIC METHOD THAT INITIALIZE TO STD VALUE THE STATIC FIELDS (THE ARE Ø BY DEFAULT BEFORE THE USE; INIT, FOR EXAMPLE)

# Initialization

- <clinit> initialization method is invoked on classes and interfaces to initialize class variables
- static initializers are executed ————— AFTER INT, DEFINED BY USER
- direct superclass need to be initialized prior
- happens on direct use: method invocation, construction, field access
- synchronized initializations: state in Class object
- <init>: initialization method for instances
  - invokespecial instruction
  - can be invoked only on uninitialized instances

KINDA CLINT  
BUT FOR INSTANCES

INITIALIZERS MODIFY STATIC FIELDS (SHARED IN THE CLASS), SO IN THE CASE OF SHARED SUPERCLASSES (E.G. OBJECT) WE HAVE TO SIGNIFY (MUTUAL EXCLUSION, LOCKS) THE NEEDS OF DIFFERENT INITIALIZED

# Initialization example (1)

```
class Super {  
    static { System.out.print("Super ");}  
}  
class One {  
    static { System.out.print("One ");}  
}  
class Two extends Super {  
    static { System.out.print("Two ");}  
}  
class Test {  
    public static void main(String[] args) {  
        One o = null;  
        Two t = new Two();  
        System.out.println((Object)o == (Object)t);  
    }  
}
```



What does **java Test** print?

Super Two False

# Initialization example (2)

```
class Super { static int taxi = 1729; }
}
class Sub extends Super {
    static { System.out.print("Sub ");}
}
class Test {
    public static void main(String[] args) {
        System.out.println(Sub.taxi);
}}
```

What does **java Test** print?

Only prints "1729" 

A reference to a static field (§8.3.1.1) causes initialization of only the class or interface that actually declares it, even though it might be referred to through the name of a subclass, a subinterface, or a class that implements an interface. (page 385 of [JLS-8])

# Finalization

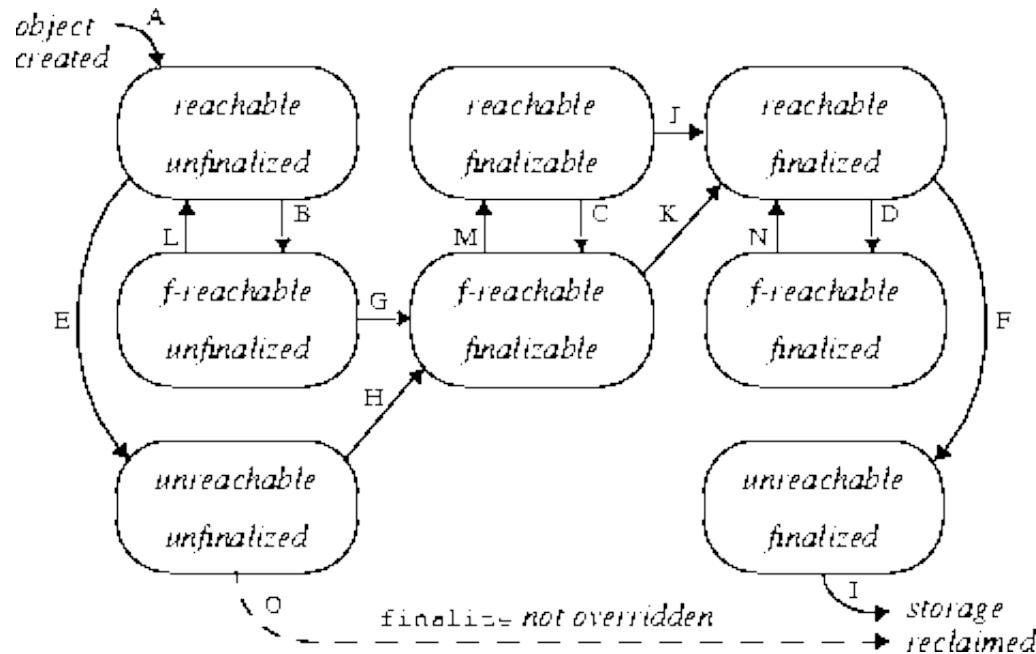
- Invoked just before garbage collection 
- JLS does not specify when it is invoked 
- Also does not specify which thread
- No automatic invocation of super's finalizers 
- Very tricky! 

```
void finalize() {  
    classVariable = this; // the object is reachable again  
}
```

- Each object can be
  - Reachable, finalizer-reachable, unreachable 
  - Unfinalized, finalizable, finalized

# Finalization State Diagram

<https://notendur.hi.is/snorri/SDK-docs/lang/lang083.htm>



**finalize()** is never called a second time on the same object, but it can be invoked as any other method!

# JVM Exit



- `classFinalize` similar to object finalization
- A class can be unloaded when
  - no instances exist
  - class object is unreachable
- JVM exits when:
  - all its non-daemon threads terminate
  - `Runtime.exit` or `System.exit` assuming it is secure
- finalizers can be optionally invoked on all objects just before exit